

CC Lecture 15

Prepared for: 7th Sem, CE, DDU

Prepared by: Niyati J. Buch

Basic abstraction of **trace based algorithm**

(e.g mark and sweep)

- All **trace-based algorithms** compute the set of **reachable** objects and then take the **complement** of this set.
- Memory is therefore recycled as follows:
 - a) The program or mutator runs and **makes allocation** requests.
 - b) The garbage collector **discovers reachability** by tracing.
 - c) The garbage collector **reclaims the storage** for unreachable objects.

Four states for chunks of memory

1. Free state

- A chunk is in the Free state if it is ready to be allocated.

2. Unreached state

- A chunk is in the Unreached state at any point during garbage collection if its reachability has not yet been established.

3. Un-scanned state

- A chunk is in the Un-scanned state if it is known to be reachable, but its pointers have not yet been scanned.


4. Scanned state

- Every Un-scanned object will eventually be scanned and transition to the Scanned state.

Scanned state

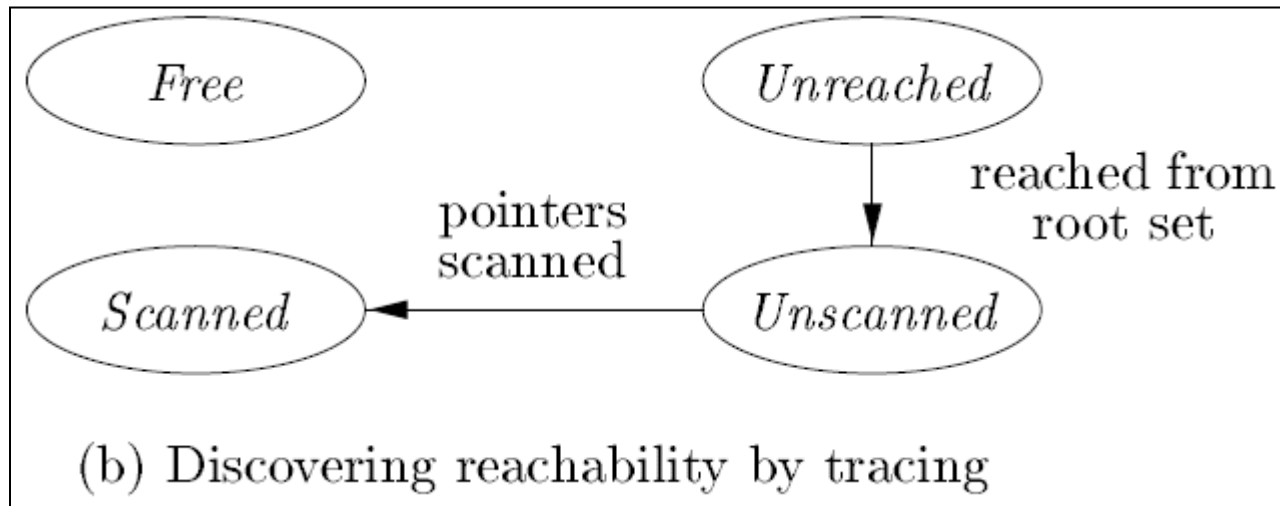
- Every **Un-scanned** object will eventually be scanned and transition to the **Scanned** state.
- To **scan** an object, we **examine** each of the pointers within it and **follow** those **pointers** to the objects to which they refer.
- If a reference is to an **Unreached** object, then that object is put in the **Un-scanned** state.
- When the scan of an object is completed, that object is placed in the **Scanned** state.
- A **Scanned** object can only contain references to other **Scanned** or **Un-scanned** objects, and never to **Unreached** objects.

States of memory in a garbage collection cycle




- Whenever a chunk is allocated by the memory manager, its state is set to **Unreached** as in Fig. (a)

States of memory in a garbage collection cycle



- The transition to **Un-scanned** from **Unreached** occurs when we discover that a chunk is **reachable** as in Fig. (b).
- When the scan of an object is completed, that object is placed in the **Scanned** state.

States of memory in a garbage collection cycle



- When no objects are left in the Un-scanned state, the computation of reachability is complete.
- Objects left in the **Unreached** state at the end are **truly unreachable**.
- The garbage collector reclaims the space they occupy and places the chunks in the **Free** state.(the solid transition in Fig. (c))
- To get ready for the next cycle of garbage collection, objects in the **Scanned** state are returned to the **Unreached** state.(the dashed transition in Fig. (c))

Mark-and-Sweep Algorithm

- **Mark-and-Sweep garbage-collection** algorithm(s) are straightforward, stop-the-world algorithm(s) that find all the unreachable objects, and put them on the list of free space.
- The algorithm has **two** phases
 - visits and “**marks**” all the reachable objects in the first tracing step
 - then “**sweeps**” the entire heap to free up unreachable objects.

Mark-and-Sweep Algorithm

- **INPUT:**
 - A **root set** of objects, a **heap**, and a free **list**, called **Free**, with **all the unallocated chunks** of the heap.
 - All chunks of space are marked with boundary tags to indicate their free/used status and size.
- **OUTPUT:**
 - A modified **Free** list after **all the garbage** has been **removed**.

Mark-and-Sweep Algorithm

- The algorithm uses several simple data structures.
 - List **Free** holds objects known to be free.
 - A list called **Unscanned**, holds objects that we have determined are reached, but whose successors(other objects can be reached through them) have not yet been considered.
 - The **Unscanned** list is empty initially.
 - Additionally, each object includes a **bit** to indicate whether it has been reached(the **reached-bit**).
 - Before the algorithm begins, all allocated objects have the **reached-bit** set to **0**.

Mark-and-Sweep Algorithm - Mark

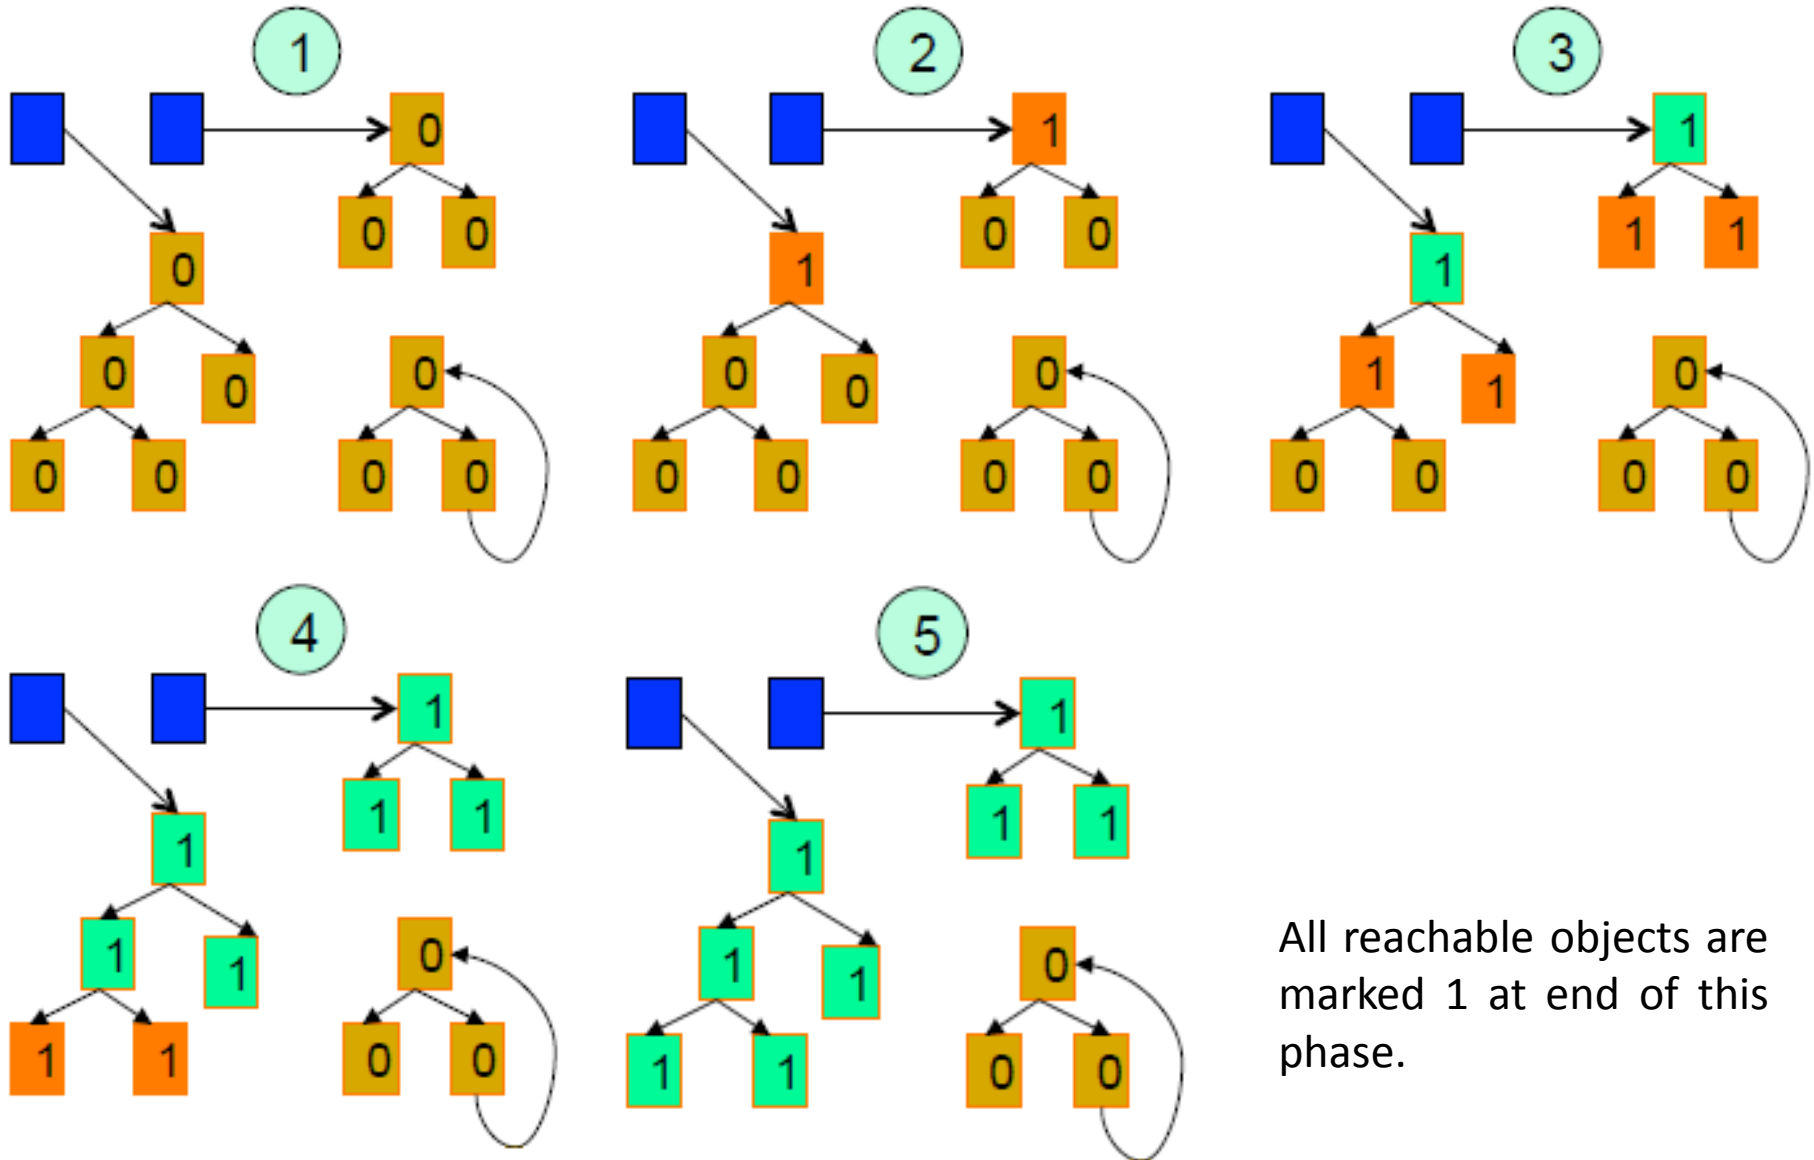
```
/* marking phase */
1) add each object referenced by the root set to list Unscanned
   and set its reached-bit to 1;
2) while (Unscanned  $\neq \emptyset$ ) {
3)   remove some object o from Unscanned;
4)   for (each object o' referenced in o) {
5)     if (o' is unreachable; i.e., its reached-bit is 0) {
6)       set the reached-bit of o' to 1;
7)       put o' in Unscanned;
           }
       }
   }
}
```

Mark-and-Sweep Algorithm - Sweep

```
/* sweeping phase */
8)  Free =  $\emptyset$ ;
9)  for (each chunk of memory o in the heap) {
10)      if (o is unreachable, i.e., its reached-bit is 0) add o to Free;
11)      else set the reached-bit of o to 0;
    }
```


Mark-and-Sweep – Example

(NPTEL Course on Principles of Compiler Design by Y.N. Srikant)



Mark-and-Sweep – Example

(NPTEL Course on Principles of Compiler Design by Y.N. Srikant)



Optimizing Mark-and-Sweep

- The final step in the basic mark-and-sweep algorithm is **expensive** because there is no easy way to find only the unreachable objects without **examining the entire heap**.
- So the time taken is proportional to the size of the heap.
- An improved algorithm, by **Baker**, keeps a list of all allocated objects.
- To find the set of unreachable objects, which we must return to free space, we take the set difference of the allocated objects and the reached objects.

Baker's Mark-and-Sweep collector

- **INPUT:**
 - A **root set** of objects, a **heap**, a free list **Free**, and a list of allocated objects, which we refer to as **Unreached**.
- **OUTPUT:**
 - Modified lists **Free** and **Unreached**, which holds allocated objects.

Baker's Mark-and-Sweep collector

- 1) $Scanned = Unscanned = \emptyset$;
- 2) move objects referenced by the root set from *Unreached* to *Unscanned*;
- 3) **while** ($Unscanned \neq \emptyset$) {
- 4) move object o from *Unscanned* to *Scanned*;
- 5) **for** (each object o' referenced in o) {
- 6) **if** (o' is in *Unreached*)
- 7) move o' from *Unreached* to *Unscanned*;
- 8) }
- 9) }
- 10) $Free = Free \cup Unreached$;
- 11) $Unreached = Scanned$;

Relocating Collectors

- **Relocating collectors** move reachable objects around in the heap to **eliminate memory fragmentation**.
- It is common that the space occupied by reachable objects is much smaller than the freed space.
- Instead of freeing the holes individually, relocate all the reachable objects into one end of the heap, leaving the entire rest of the heap as one free chunk.
- As GC already analyzed every reference within the reachable objects
- So this and references in root set is required to be changed.

Advantages

- Having all the reachable objects in contiguous locations **reduces fragmentation** of the memory space.
- Also, by making the data occupy fewer cache lines and pages, relocation **improves** a program's **temporal and spatial locality**, since new objects created at about the same time are allocated nearby chunks.
- Objects in nearby chunks can benefit from **prefetching** if they are used together.
- Further, the **data structure** for maintaining free space is **simplified**; instead of a free list, all we need is a **pointer** free to the beginning of the one free block.

Types of Relocating Collectors

- **Relocating collectors** vary in whether they relocate in place or reserve space ahead of time for the relocation:
 1. A **Mark-and-Compact collector**, described in this section, compacts objects in place.
 - Relocating in place reduces memory usage.
 2. The more efficient and popular **Copying Collector** moves objects from one region of memory to another.
 - Reserving extra space for relocation allows reachable objects to be moved as they are discovered.

3 phases of Mark-and-Compact collector

1. First is a **marking** phase, similar to that of the mark-and-sweep algorithms described previously.
2. Second, the algorithm **scans** the allocated section of the heap and **computes a new address** for each of the reachable objects.
 - New addresses are assigned from the **low end** of the heap, so there are no holes between reachable objects.
 - The new address for each object is recorded in a structure called **NewLocation**.
3. Finally, the algorithm **copies** objects to their new locations, updating all references in the objects to point to the corresponding new locations.
 - The needed addresses are found in **NewLocation**.